

Color and Blending

Prerequisites

None; this module discusses color issues from first principles.

Introduction

Color is a fundamental concept for computer graphics. We need to be able to define colors for our graphics that represent good approximations of real-world colors, and we need to be able to manipulate colors as we develop our applications.

There are many ways to specify colors, but all depend principally on the fact that the human visual system generally responds to colors through the use of three kinds of cells in the retina of the eye. This response is complex and includes both physical and psychological processes, but the fundamental fact of three kinds of stimulus is maintained by all the color models in computer graphics. For most work, the usual model is the RGB (Red, Green, Blue) color model that matches in software the physical design of computer monitors, which are made with a pattern of three kinds of phosphor which emit red, green, and blue light when they are excited by an electron beam. This RGB model is used for color specification in almost all computer graphics APIs, and it is the basis for the discussion here. There are a number of other models of color, and we refer you to textbooks and other sources, especially Foley et al. [1], for additional discussions on color models and for more complete information on converting color representations from one model to another.

Because the computer monitor uses three kinds of phosphor, and each phosphor emits light levels based on the energy of the electron beam that is directed at it, a common approach is to specify a color by the level of each of the three primaries. These levels are a proportion of the maximum light energy possible for that primary, so a color is specified by a triple (r, g, b) where each of the three components is a real number between 0.0 and 1.0, inclusive. The number represents the proportion of the available color of that primary hue that is desired for the pixel. Thus the higher the number for a component, the brighter is the light in that color, so black is represented by $(0.0, 0.0, 0.0)$ and white by $(1.0, 1.0, 1.0)$. The RGB primaries are represented respectively by $(1.0, 0.0, 0.0)$ — red, $(0.0, 1.0, 0.0)$ — green, and $(0.0, 0.0, 1.0)$ — blue; that is, colors that are fully bright in a single primary component and totally dark in the other primaries. Other colors are a mix of the three primaries as needed.

Internally, most graphics hardware does not deal with colors using floating-point numbers. Instead integer values are used to save space and to speed operations, with the exact representation and storage of those integers depending on the number of bits per color per pixel and on other hardware design issues. This distinction sometimes comes up in considering details of OpenGL operations, but is generally something that can be ignored. The color-generation process itself is surprisingly complex, because the viewing device must generate perceptually-linear values, and most hardware generates color with exponential, not linear, properties. All these are hidden from the OpenGL programmer, however, and are managed after being translated from the floating-point representations of the colors; this gives us a device-independent representation that allows OpenGL programs to work relatively the same across a wide range of platforms.

In addition to dealing with the color of light, modern graphics systems add a fourth component to the question of color. This fourth component is called “the alpha channel” because that was its original notation [2], and it represents the opacity of the material that is being modeled. As is the case with color, this is represented by a real number between 0.0 (no opacity — completely transparent) and 1.0 (completely opaque — no transparency). This is used to allow you to create

objects that you can see through at some level, and can be a very valuable tool when you want to be able to see more than just the things at the front of a scene. However, transparency is not determined globally by the graphics API; it is determined by compositing the new object with whatever is already present in the Z-buffer. Thus if you want to create an image that contains many levels of transparency, you will need to pay careful attention to the sequence in which you draw your objects, drawing the furthest first in order to get correct attenuation of the colors of background objects.

Definitions

The RGB cube

The RGB color model is associated with a geometric presentation of a color space. That space is a cube consisting of all points (r, g, b) with each of r , g , and b having a value that is a real number between 0 and 1. Because of the easy analogy between color triples and space triples, every point in the unit cube can be easily identified with a RGB triple representation of a color. This gives rise to the notion of the RGB color cube that is seen in every graphics text (and thus that we won't repeat here).

To illustrate the numeric properties of the RGB color system, we will create the edges of the color cube as shown in Figure 5.1 below, which has been rotated to illustrate the colors more fully. To do this, we create a small cube with a single color, and then draw a number of these cubes around the edge of the geometric unit cube, with each small cube having a color that matches its location. We see the origin $(0,0,0)$ corner, farthest from the viewer, mostly by its absence because of the black background, and the $(1,1,1)$ corner nearest the viewer as white. The three axis directions are the pure red, green, and blue corners. The code for this example is provided in the file `colorcubes.c`, and you go into the code and change the background color so you can see the origin by using use a shade of gray (a color whose r , g , and b components are all equal) to avoid having some other point of the cube disappear.




Figure 5.1: tracing the colors of the edges of the RGB cube

This figure suggests the nature of the RGB cube, but a view of the entire RGB cube is provided by the Figure 5.2, showing the cube from two points of view (from the white vertex and from the black vertex) so you can see the entire gamut of colors on the surface of the cube. Note that the three vertices closest to the white vertex are the cyan, magenta, and yellow vertices, while the three vertices closest to the black vertex are the red, green, and blue vertices. This illustrates the additive nature of the RGB color model, with the colors getting lighter as the amounts of the primary colors increases. This will be explored later and will be contrasted to the subtractive nature of other color models. Code to create the views of the RGB cube is provided in the `RGBspace.c` file that is included with this module. Not shown is the center diagonal of the RGB cube from $(0, 0, 0)$ to

(1, 1, 1) that corresponds to the colors with equal amounts of each primary; these are the gray colors that provide neutral backgrounds as suggested above.




Figure 5.2: two views of the RGB cube — from white (left) and black (right)

Color is ubiquitous in computer graphics, and can be applied to both objects and lights. In this module we only think about the colors of objects, and save the color of lights, and the difference in the way we specify color in a lighted and unlighted scene, for a later module on lighting. In general, the behavior of a scene will reflect both these attributes — if you have a red object and illuminate it with a blue light, your object will seem to be black, because a red object reflects no blue light and the light contains no other color than blue.

Other color models

There are times when the RGB model is not easy to use. Few of us think of a particular color in terms of the proportions of red, green, and blue that are needed to create it, so there are other ways to think about color that make this more intuitive. And there are some processes for which the RGB approach does not model the reality of color production. So we need to have a wider range of ways to model color to accomodate these realities.

A more intuitive approach to color is found with either of the HSV (Hue-Saturation-Value) or HLS (Hue-Lightness-Saturation) models. These models represent color as a hue (intuitively, a descriptive variation on a standard color such as red, or magenta, or blue, or cyan, or green, or yellow) that is modified by setting its value (a property of darkness or lightness) and its saturation (a property of brightness). This lets us find numerical ways to say "the color should be a dark, vivid reddish-orange" by using a hue that is to the red side of yellow, has a relatively low value, and has a high saturation.

Just as there is a geometric model for RGB color space, there is one for HSV color space: a cone with a flat top, as shown in Figure 5.3 below. The distance around the circle in degrees represents the hue, starting with red at 0, moving to green at 120, and blue at 240. The distance from the vertical axis to the outside edge represents the saturation, or the amount of the primary colors in the particular color. This varies from 0 at the center (no saturation, which makes no real coloring) to 1 at the edge (fully saturated colors). The vertical axis represents the value, from 0 at the bottom (no color, or black) to 1 at the top. So a HSV color is a triple representing a point in or on the cone, and the "dark, vivid reddish-orange" color would be something like (40.0, 1.0, 0.7). Code to display this geometry interactively is included in the file `HSVspace.c`, and the interactive nature of the program gives a much better view of the space.

The shape of the HSV space can be a bit confusing. The top surface represents all the lighter colors based on the primaries, because colors getting lighter have the same behavior as colors getting less saturated. The reason the geometric model tapers to a point at the bottom is that there is no real color variation near black. In this model, the gray colors are the colors with a saturation of

0, which form the vertical center line of the cone. For such colors, the hue is meaningless, but it still must be included.



Figure 5.3: three views of HSV color space: side (left), top (middle), bottom (right)

In the HLS color model, shown in Figure 5.4, the geometry is much the same as the HSV model but the top surface is stretched into a second cone. Hue and saturation have the same meaning as HSV but lightness replaces value, and lightness corresponds to the brightest colors at a value of 0.5. The rationale for the dual cone that tapers to a point at the top as well as the bottom is that as colors get lighter, they lose their distinctions of hue and saturation in a way that is very analogous with the way colors behave as they get darker. In some ways, the HLS model seems to come closer to the way people talk about “tints” and “tones” when they talk about paints, with the strongest colors at lightness 0.5 and becoming lighter (tints) as the lightness is increased towards 1.0, and becoming darker (tones) as the lightness is decreased towards 0.0. Just as in the HSV case above, the grays form the center line of the cone with saturation 0, and the hue is meaningless.




Figure 5.4: the HLS double cone from the red (left), green (middle), and blue(right) directions.

The top and bottom views of the HLS double cone look just like those of the HSV single cone, but the side views of the HLS double cone are quite different. Figure 5.4 shows the HLS double cone from the three primary-color sides: red, green, and blue respectively. The views from the top or bottom are exactly those of the HSV cone and so are now shown here. The images in the figure do not show the geometric shape very well; the code in `HLSspace.c` allows you to interact with the model and see it more effectively in 3-space.

There are relatively simple functions that convert a color defined in one space into the same color as defined in another space. We do not include all these functions in these notes, but they are covered in [1], and the functions to convert HSV to RGB and to convert HLS to RGB are included in the sample code that produced these figures.

All the color models above are based on colors presented on a computer monitor or other device where light is emitted to the eye. Such colors are called *emissive* colors, and operate by adding light at different wavelengths as different screen cells emit light. The fact that most color presented by programming comes from a screen makes this the primary way we think about color in computer graphics systems. This is not the only way that color is presented to us, however. When you read these pages in print, and not on a screen, the colors you see are generated by light that is reflected from the paper through the inks on the page. Such colors can be called *transmissive* colors and operate by subtracting colors from the light being reflected from the page. This is a totally different process and needs separate treatment. Figure 5.5 below illustrates this principle. The way the RGB add to produce CMY and eventually white shows why emissive colors are sometimes called additive colors, while the way CMY produce RGB and eventually black shows why transmissive colors are sometimes called subtractive colors.



Figure 5.5: emissive colors (left) and transmissive colors (right)

Transmissive color processes use inks or films that transmit only certain colors while filtering out all others. Two examples are the primary inks for printing and the films for theater lights; the



Figure 5.6: color separations for printing

primary values for transmissive color are cyan (which transmits both blue and green), magenta (which transmits both blue and red), and yellow (which transmits both red and green). In principle, if you use all three inks or filters (cyan, magenta, and yellow), you should have no light transmitted and so you should see only black. In practice, actual materials are not perfect and allow a little off-color light to pass, so this would produce a dark and muddy gray (the thing that printers call “process black”) so you need to add an extra “real” black to the parts that are intended to be really black. This cyan-magenta-yellow-black model is called CMYK color and is the basis for printing and other transmissive processes. It is used to create color separations that combine to form full-color images as shown in Figure 5.6, which shows a full-color image (left) and the sets of yellow, cyan, black, and magenta separations (right-hand side, clockwise from top left) that are used to create plates to print the color image. We will not consider the CMYK model further in this discussion because its use is in printing and similar technologies, but not in graphics programming.

Color blending with the alpha channel

A color can be represented in OpenGL as more than just a RGB triple; it can also include a blending level (sometimes thought of as a transparency level) so that anything with this color will have a color blending property. Thus color is also represented by a quadruple (r,g,b,a). This transparency level *a* is called the alpha value, and its value is a number between 0.0 and 1.0 that is actually a measure of opacity instead of transparency. That is, if you use standard kinds of blending functions and if the alpha value is 1.0, the color is completely opaque, but in the same situation if the alpha value is 0.0, the color is completely transparent. However, we are using the term “transparent” loosely here, because the real property represented by the alpha channel is *blending*, not transparency. The alpha channel was invented to permit image image compositing [2] in which an image could be laid over another image and have part of the underlying image show through. So while we may say “transparent” we really mean blended.

This difference between blended and transparent colors can be very significant. If we think of transparent colors, we are modeling the logical equivalent of colored glass. This kind of material embodies transmissive, not emissive, colors — only certain wavelengths are passed through, while the rest are absorbed. But this is not the model OpenGL uses for the alpha value; blended colors operate by averaging emissive RGB colors, which is the opposite of the transmissive model implied by transparency. The difference can be important in creating the effects you need in an image. There is an additional issue to blending because averaging colors in RGB space may not result in the intermediate colors you would expect; the RGB color model is one of the worse color models for perceptual blending but we have no choice in OpenGL.

Enabling blending

In order to use transparent colors, you must specify that you want the blending enabled and you must identify the way the color of the object you are drawing will be blended with the color that has already been defined. This is done with two simple function calls:

```
glEnable(GL_BLEND);  
glBlendFunc(GL_SRC_ALPHA, GL_ONE_MINUS_SRC_ALPHA);
```

The first is a case of the general enabling concept for OpenGL; the system has many possible capabilities and you can select those you want by enabling them. This allows your program to be more efficient by keeping it from having to carry out all the possible operations in the rendering pipeline. The second allows you to specify how you want the color of the object you are drawing to be blended with the color that has already been specified for each pixel. If you use this blending function and your object has an alpha value of 0.7, for example, then the color of a pixel after it has been drawn for your object would be 70% the color of the new object and 30% the color of whatever had been drawn up to that point.

There are many options for the OpenGL blending function. The one above is the most commonly used and simply computes a weighted average of the foreground and background colors, where the weight is the alpha value of the foreground color. In general, the format for the blending specification is

```
glBlendFunc(src, dest)
```

and there are many symbolic options for the source (src) and destination (dest) blending values; the OpenGL manual covers them all.

Blending Effects

Blending creates some significant challenges if we want to create the impression of transparency. To begin, we make the simple observation that if something is intended to seem transparent to some degree, you must be able to see things behind it. This suggests a simple first step: if you are working with objects having their alpha color component less than 1.0, it is useful and probably important to allow the drawing of things that might be behind these objects. To do that, you should draw all solid objects (objects with alpha component equal to 1.0) before drawing the things you want to seem transparent, and then turn off the depth buffer by disabling the depth test with `glDisable(GL_DEPTH_TEST)` before drawing items with blended colors, and enabling the depth test again after drawing them. This at least allows the possibility that some concept of transparency is allowed.

But it may not be enough to do this, and in fact this attempt at transparency may lead to more confusing images than leaving the depth test intact. Let us consider the case that you have three objects to draw, and that you will be drawing them in a particular order. For the discussion, let's assume that the objects are numbered 1, 2, and 3, that they have colors C1, C2, and C3, that you draw them in the sequence 1, 2, and 3, that they line up from the eye but have a totally white background behind them, and that each color has $\alpha = 0.5$. Let's assume further that we are not using the depth buffer so that the physical ordering of the objects is not important. The layout is shown in Figure 5.7. And finally, let's further assume that we've specified the blend functions as suggested above, and consider the color that will be drawn to the screen where these objects lie.




Figure 5.7: the sequence for drawing the objects

When we draw the first object, the frame buffer will have color C1; no other coloring is involved. When we draw the second object on top of the first, the frame buffer will have color $0.5 \cdot C1 + 0.5 \cdot C2$, because the foreground (C2) has alpha 0.5 and the background (C1) is included with weight $0.5 = 1 - 0.5$. Finally, when the third object is drawn on top of the others, the color will be $0.5 \cdot C3 + 0.5 \cdot (0.5 \cdot C1 + 0.5 \cdot C2)$, or $0.5 \cdot C3 + 0.25 \cdot C2 + 0.25 \cdot C1$. That is, the color of the most recent object drawn is emphasized much more than the color of the other objects. This shows up clearly in the right-hand part of Figure 5.8 below, where the red square is drawn after the other two squares. On the other hand, if you had drawn object three before object 2, and object 2 before object 1, the color would have been $0.5 \cdot C1 + 0.25 \cdot C2 + 0.25 \cdot C3$, so the order in which you draw things, not the order in which they are placed in space, determines the color.

But this again emphasizes a difference between blending and transparency. If we were genuinely modeling transparency, it would not make any difference which object were placed first and which last; each would subtract light in a way that is independent of its order. So this represents another challenge if you would want to create an illusion of transparency with more than one non-solid object.

Some examples

Example: An object with partially transparent faces

If you were to draw a piece of the standard coordinate planes and to use colors with alpha less than 1.0 for the planes, you would be able to see through each coordinate plane to the other planes as though the planes were made of some partially-transparent plastic. We have modeled a set of three squares, each lying in a coordinate plane and centered at the origin, and each defined as having a rather low alpha value of 0.5 so that the other squares are supposed to show through. In this section we consider the effects of a few different drawing options on this view.

The left-hand side of Figure 5.8 below shows the image we get with these colors when we leave the depth test active. Here what you can actually “see through” depends on the order in which you draw the objects. With the depth test enabled, the presence of a transparent object close to your eye prevents the writing of its blend with an object farther away. Because of this, the first coordinate plane you draw is completely opaque to the other planes, even though we specified it as being partly transparent, and a second coordinate plane allows you to see through it to the first plane but is fully opaque to the second plane. We drew the blue plane first, and it is transparent only to the background (that is, it is darker than it would be because the black background shows through). The green plane was drawn second, and that only allows the blue plane to show through. The red plane was drawn third, and it appears to be fully transparent and show through to both other planes. In the actual working example, you can use keypresses to rotate the planes in space; note that as you do, the squares you see have the same transparency properties in any position.




Figure 5.8: the partially transparent coordinates planes (left); the same coordinate planes fully transparent but with same alpha (center); the same coordinate planes with adjusted alpha (right)

However, in the image in the center of Figure 5.8, we have disabled the depth test, and this presents a more problematic situation. In this case, the result is something much more like transparent planes, but the transparency is very confusing because the last plane drawn, the red plane, always seems to be on top because its color is the brightest. This figure shows that the OpenGL attempt at transparency is not necessarily a desirable property; it is quite difficult to get information about the relationship of the planes from this image. Thus one would want to be careful with the images one would create whenever one chose to work with transparent or blended

images. Because this figure is created by exactly the same code as the one above with `GL_BLEND` disabled instead of enabled, the code for the example is not included here.

Finally, we change the alpha values of the three squares to account for the difference between the weights in the final three-color section. Here we use 1.0 for the first color (blue), 0.5 for the second color (green) but only 0.33 for red, and we see that this final image, the right-hand image in Figure 5.8, has the following color weights in its various regions:

- 0.33 for each of the colors in the shared region,
- 0.5 for each of blue and green in the region they share,
- 0.33 each for red and green in the region they share,
- 0.33 for red and 0.67 for blue in the region they share,
- the original alpha values for the regions where there is only one color.

Note that the “original alpha values” gives us a solid blue, a fairly strong green, and a weak red as stand-alone colors. This gives us a closer approximation to the appearance actual transparency for these three colors, with a particular attention to the clear gray in the area they all cover, but there are still some areas that don’t quite work. To get even this close, however, we must analyze the rendering carefully and we still cannot quite get a perfect appearance. The code for all these examples is in the file `alphacolor.c`, which is presented with blending disabled and the adjusted values of alpha for the right-hand image in the figure.

A word to the wise...

Recall that the alpha value represents a blending proportion, not transparency. This blending is applied by comparing the color of an object with the current color of the image at a pixel, and coloring the pixel by blending the current color and the new color according to one of several rules that you can choose with `glBlendFunc(...)` as noted above. This capability allows you to build up an image by blending the colors of parts in the order in which they are rendered, and we saw that the results can be quite different if objects are received in a different order. Blending does not treat parts as being transparent, and so there are some images where OpenGL simply does not blend colors in the way you expect from the concept.

If you really do want to try to achieve the illusion of full transparency, you are going to have to do some extra work. You will need to be sure that you draw the items in your image starting with the item at the very back and proceeding to the frontmost item. This process is sometimes called Z-sorting and it can be tricky because objects can overlap or the sequence of objects in space can change as you apply various transformations to the scene. In the example above, the squares actually intersect each other and could only be sorted if each were broken down into four separate sub-squares. And even if you can get the objects sorted once, the order would change if you rotated the overall image in 3-space, so you would possibly have to re-sort the objects after each rotation. In other words, this would be difficult.

As always, when you use color you must consider carefully the information it is to convey. Color is critical to convey the relation between a synthetic image and the real thing the image is to portray, of course, but it can be used in many more ways. One of the most important is to convey the value of some property associated with the image itself. As an example, the image can be of some kind of space (such as interstellar space) and the color can be the value of something that occupies that space or happens in that space (such as jets of gas emitted from astronomical objects where the value is the speed or temperature of that gas). Or the image can be a surface such as an airfoil (an airplane wing) and the color can be the air pressure at each point on that airfoil. Color can even be used for displays in a way that carries no meaning in itself but is used to support the presentation, as in the Chromadepth™ display we will discuss below in the texture mapping module. But never use color without understanding the way it will further the message you intend in your image.

Code examples

Example: A model with parts having a full spectrum of colors

The code in the file `cubecolors.c` uses a technique you have not yet seen to create all the small cubes that make up the visible edges of the RGB cube. In this code, we use only a simple cube we defined ourselves (not that it was too difficult!) to draw each cube, setting its color by its location in the space:

```
typedef GLfloat color [4];
color cubecolor;

cubecolor[0] = r; cubecolor[1] = g; cubecolor[2] = b;
cubecolor[3] = 1.0;
glColor4fv(cubecolor);
```

However, we don't change the geometry of the cube before we draw it, so how does this cube move all over the space? In fact, we only use the cube we defined, which is a cube whose sides all have length two and which is centered on the origin. So how do we get the small cubes we actually draw? Our technique is to use *transformations* to define the size and location of each of the cubes, with *scaling* to define the size of the cube and *translation* to define the position of the cube, as follows:

```
glPushMatrix();
glScalef(scale, scale, scale);
glTranslatef(-SIZE+(float)i*2.0*scale*SIZE, SIZE, SIZE);
cube((float)i/(float)NUMSTEPS, 1.0, 1.0);
glPopMatrix();
```

Note that we include a technique of pushing the current modeling transformation onto the current transformation stack, applying the translation and scaling transformations to the transformation in use, drawing the cube, and then popping the current transformation stack to restore the previous modeling transformation. This is discussed further in a later module where we discuss the general capabilities of transformations.

Example: the HSV cone

There are two functions of interest here. The first is the conversion from HSV colors to RGB colors; this is taken from [1] as indicated, and is based upon a geometric relationship between the cone and the cube, which is much clearer if you look at the cube along a diagonal between two opposite vertices. The second function does the actual drawing of the cone with colors generally defined in HSV and converted to RGB for display, and with color smoothing handling most of the problem of shading the cone. The full code for this program is included with this module.

```
void convertHSV2RGB(float h, float s, float v, float *r, float *g, float *b)
{
    // conversion from Foley et.al., fig. 13.34, p. 593
    float f, p, q, t;
    int k;

    if (s == 0.0) { // achromatic case
        *r = *g = *b = v;
    }
    else { // chromatic case
        if (h == 360.0) h=0.0;
        h = h/60.0;
        k = (int)h;
        f = h - (float)k;
        p = v * (1.0 - s);
        q = v * (1.0 - s * f);
        t = v * (1.0 - s * (1.0 - f));
        *r = (k==0)?p:(k==1)?q:(k==2)?t:p;
        *g = (k==0)?t:(k==1)?p:(k==2)?q:t;
        *b = (k==0)?q:(k==1)?t:(k==2)?p:q;
    }
}
```

```

        q = v * (1.0 - (s * f));
        t = v * (1.0 - (s * (1.0 - f)));
        switch (k) {
            case 0: *r = v; *g = t; *b = p; break;
            case 1: *r = q; *g = v; *b = p; break;
            case 2: *r = p; *g = v; *b = t; break;
            case 3: *r = p; *g = q; *b = v; break;
            case 4: *r = t; *g = p; *b = v; break;
            case 5: *r = v; *g = p; *b = q; break;
        }
    }
}

void HSV(void)
{
#define NSTEPS 36
#define steps (float)NSTEPS
#define TWOPI 6.28318

    int i;
    float r, g, b;

    glBegin(GL_TRIANGLE_FAN);        //      cone of the HSV space
        glColor3f(0.0, 0.0, 0.0);
        glVertex3f(0.0, 0.0, -2.0);
        for (i=0; i<=NSTEPS; i++) {
            convert(360.0*(float)i/steps, 1.0, 1.0, &r, &g, &b);
            glColor3f(r, g, b);
            glVertex3f(2.0*cos(TWOPI*(float)i/steps),
                      2.0*sin(TWOPI*(float)i/steps),2.0);
        }
    glEnd();
    glBegin(GL_TRIANGLE_FAN);        //      top plane of the HSV space
        glColor3f(1.0, 1.0, 1.0);
        glVertex3f(0.0, 0.0, 2.0);
        for (i=0; i<=NSTEPS; i++) {
            convert(360.0*(float)i/steps, 1.0, 1.0, &r, &g, &b);
            glColor3f(r, g, b);
            glVertex3f(2.0*cos(TWOPI*(float)i/steps),
                      2.0*sin(TWOPI*(float)i/steps),2.0);
        }
    glEnd();
}

```

Example: the HLS double cone

The conversion itself takes two functions, while the function to display the double cone is so close to that for the HSV model that we do not include it here. The source of the conversion functions is again Foley et al.

```

void convertHLS2RGB(float h,float l,float s,float *r,float *g,float *b)
{
    // conversion from Foley et.al., Figure 13.37, page 596
    float m1, m2;

    if (l <= 0.5) m2 = l*(1.0+s);
    else          m2 = l + s - l*s;
    m1 = 2.0*l - m2;

```

```

    if (s == 0.0) {    // achromatic cast
        *r = *g = *b = 1;
    }
    else {    // chromatic case
        *r = value(m1, m2, h+120.0);
        *g = value(m1, m2, h);
        *b = value(m1, m2, h-120.0);
    }
}

float value( float n1, float n2, float hue)
{    // helper function for the HLS->RGB conversion
    if (hue > 360.0) hue -= 360.0;
    if (hue < 0.0)    hue += 360.0;
    if (hue < 60.0)    return( n1 + (n2 - n1)*hue/60.0 );
    if (hue < 180.0) return( n2 );
    if (hue < 240.0) return( n1 + (n2 - n1)*(240.0 - hue)/60.0 );
    return( n1 );
}

```

Example: An object with partially transparent faces

The code in the file `alphacolor.c` has a few properties that are worth noting. This code draws three squares in space, each centered at the origin and lying within one of the coordinate planes. These three squares are colored according to the following declaration

```

GLfloat color0[]={1.0, 0.0, 0.0, 0.5}, // R
          color1[]={0.0, 1.0, 0.0, 0.5}, // G
          color2[]={0.0, 0.0, 1.0, 0.5}; // B

```

These colors are the full red, green, and blue colors with a 0.5 alpha value, so when each square is drawn it uses 50% of the background color and 50% of the square color. You will see that blending in Figure 5.8 for this example.

As we saw in the example above, the color of each part is specified just as the part is drawn:

```

point3 plane0[4]={ {-1.0, 0.0, -1.0}, // X-Z plane
                   {-1.0, 0.0,  1.0},
                   { 1.0, 0.0,  1.0},
                   { 1.0, 0.0, -1.0} };

```

This is not necessary if many of the parts had the same color; once a color is specified, it is used for anything that is drawn until the color is changed.

Example code

The following source files are examples from this module.

- `alphacolor.c`
- `colorcubes.c`
- `RGBspace.c`
- `HSVspace.c`: defining colors in HSV space and displaying them in RGB, including displaying the HSV cone itself
- `HLSspace.c`: defining colors in HLS space and displaying them in RGB, including displaying the HLS double cone itself

Science projects

- Create objects having a property that can be encoded with color, and display the object using color to demonstrate that property. For example, consider the problem of heat transport in a bar, with some points of the bar being held at different temperatures by heat sources or sinks,

and draw the bar with colors illustrating the temperatures at different parts of the bar. (Note that there are conventions about colors representing heat — red represents heat, and a “cooler” color such as blue or green (or cyan, which is (0,1,1)) representing cold.)

- There are some fields in which colors have standardized meanings. In chemistry, there are standard colors for some atoms in molecular displays, for example. Identify some such fields and examine how the colors are used to show structures or other properties in this case.
- If you do not have lighting and shading, it can be difficult to see shapes in graphic displays. It is possible to display shapes by using color to give hints to the eye. For example, topographic maps use colored lines at constant elevations to tell us where a landscape is largely flat (the lines are far apart) or very steep (the lines are very close together). If we are drawing a surface and have defined the surface in terms of a fine grid in the domain, it can be very helpful if we use color to describe height by using different colors for different heights. This allows our eyes to identify shapes much as they did for topographic maps.

References

- [1] Foley, van Dam, Feiner, and Hughes, *Computer Graphics: Principles and Practice*, 2nd edition, Addison-Wesley, 1990
- [2] Porter and Duff, “Compositing Digital Images,” *Proceedings, SIGGRAPH 84 Conference*, ACM SIGGRAPH, 1984